

Key performance indicators

Breno Gomes
Senior Solution Engineer

January 2022

Why it is important

Modern software has a large volume of interconnected components. In an ideal world the number and types of failures grow proportionally. Strong adoption of cloud functions and containers increases the complexity.

Distributed systems are updated regularly. Every change can create a new type of failure, instability or more “unknown unknowns”.

Key performance indicators are more than numbers you report during the software development lifecycle.

Metrics provide targets for teams to set goals, milestones to gauge progress, and insights that help people across the organization make informed decisions.

A good KPI, by definition, should be measurable and trackable.

KPI in the context of feature flags and experimentation

Existing and prospective customers normally go through a proof of value (PoV) before adopting or subscribing new products.

It is important to define pilot project to evaluate the feasibility of an implementation before it is widely incorporated in the software development lifecycle. It defines the processes to simplify a “minimum viable product”. A proof of value targets the financial and organisational benefits of such implementation.

Key performance indicators are fundamental to define the success criteria. They are used to compare metrics before, during and after the PoV.

Engineers are often caught off guard when they are asked to share the KPI for a PoV. That is perfectly natural unless they were briefed beforehand. This guide intends to share metrics frequently adopted during product evaluation. It is neither prescriptive nor exhaustive. Each company is unique, having freedom to choose what suits it best.

Use cases

Synchronously toggle features across platforms.

Moving from an old service to new service and using LaunchDarkly to incrementally accept traffic.

Migrating between databases by reading/writing to different datastores without multiple deployments.

Control the rollout of expensive operations during application updates.

Switching between different UI themes for given sets of client users, allowing controlled testing before release to the store.

Safely beta/alpha test features in production releases to gather feedback.

Flag for additional content to test click-through rate.

DevOps

Builds

Commits

Deployment frequency

Approving a feature release

Lead time for changes

Change volume

Change failure rate

Defect escape rate

Mean time to detect

Mean time to identify

Mean time to restore
service

Percentage of code covered
by automated tests

Application usage and
traffic

Application availability

Support tickets

Business

Marketing conversion funnel

A/B tests

Cart abandonment

Unique users

Cart optimisation

Number of sessions

Revenue per customer

Page views

Customer retention rate

Session duration

Profit margin

Geography

Technical

Perceived page load time or
mobile interaction

Web browser

Error rate: frontend, backend and
infrastructure

Mobile device

API timeout

Mobile crash rate

API error

DOM processing

API latency

Page rendering

Web and Infrastructure

Availability

Throughput

Application response time

Database execution time

Error rate

Memory footprint

CPU utilisation

Network throughput

DORA

Aspect of Software Delivery Performance*	Elite	High	Medium	Low
Deployment frequency For the primary application or service you work on, how often does your organization deploy code to production or release it to end users?	On-demand (multiple deploys per day)	Between once per day and once per week	Between once per week and once per month	Between once per month and once every six months
Lead time for changes For the primary application or service you work on, what is your lead time for changes (i.e., how long does it take to go from code committed to code successfully running in production)?	Less than one day	Between one day and one week	Between one week and one month	Between one month and six months
Time to restore service For the primary application or service you work on, how long does it generally take to restore service when a service incident or a defect that impacts users occurs (e.g., unplanned outage or service impairment)?	Less than one hour	Less than one day ^a	Less than one day ^a	Between one week and one month
Change failure rate For the primary application or service you work on, what percentage of changes to production or released to users result in degraded service (e.g., lead to service impairment or service outage) and subsequently require remediation (e.g., require a hotfix, rollback, fix forward, patch)?	0-15% ^{b,c}	0-15% ^{b,d}	0-15% ^{c,d}	46-60%

Source: <https://cloud.google.com/blog/products/devops-sre/using-the-four-keys-to-measure-your-devops-performance>